

# Model Selection & Evaluation

March 25th, 2019

# Learnable Theory

## How to define learning in ML?

- Several questions arise when designing and analyzing algorithms that learn from data. Examples of questions include:
  - ① What can be learned efficiently?
  - ② What is inherently hard to learn?
  - ③ How many examples are needed to learn successfully?
  - ④ Is there a general model of learning?
- The ***Probably Approximately Correct (PAC)*** learning framework helps defines the class of learnable concepts in terms of ***number of sample points*** needed to achieve an approximate solution, ***sample complexity***, and the time and space complexity of a learning algorithm.

- Let us denote  $\mathcal{X}$  the set of all possible examples,  $\mathcal{Y}$  the set of all possible label or target values, and that  $\mathcal{Y} = \{0, 1\}$
- A concept  $c : \mathcal{X} \mapsto \mathcal{Y}$  is a mapping from  $\mathcal{X}$  to  $\mathcal{Y}$ .
- $\mathcal{C}$  is a concept class that comprises the concepts we may wish to learn
- The **learning problem** can be formulated as follows:
  - The learner considers a fixed set of all possible concepts  $\mathcal{H}$ , called *hypothesis set*, with input sample  $S = (x^1, \dots, x^p)$  draw i.i.d according to  $\mathcal{D}$  as well as the labels  $c(x_1), \dots, c(x^p)$  with  $c \in \mathcal{C}$
  - The task comprise in using the labeled sample  $S$  to select a hypothesis  $h_s \in \mathcal{H}$  that as a small **generalization error** with respect to  $c$ .
- The generalization error of a hypothesis  $h \in \mathcal{H}$  is also known as the **risk** or **true error**.

## Generalization error

Given a hypothesis  $h \in \mathcal{H}$ , a target concept  $c \in \mathcal{C}$ , and an underlying distribution  $\mathcal{D}$ , the generalization error or risk of  $h$  is defined by

$$R(h) = \mathbb{P}_{x \sim \mathcal{D}}[h(x) \neq c(x)] = \mathbb{E}_{x \sim \mathcal{D}}[\mathbf{1}_{h(x) \neq c(x)}]$$

- Since both the distribution  $\mathcal{D}$  and the target concept  $c$  is unknown, a learner cannot direct access the generalization error. It can only measure the **empirical error** of a  $h \in \mathcal{H}$  on the labeled sample  $\mathcal{S}$

## Empirical error

### Empirical error

- Given a hypothesis  $h \in \mathcal{H}$ , a target concept  $c \in \mathcal{C}$ , and a sample  $S = (x^1, \dots, x^m)$ , the empirical error or empirical risk of  $h$  is defined by

$$\hat{R}_S(h) = \frac{1}{m} \sum_{i=1}^m 1_{h(x_i) \neq c(x_i)}$$

- The **empirical error** of  $h \in \mathcal{H}$  is its average error over the sample  $S$ , while the **generalization error** is its expected error based on the distribution  $\mathcal{D}$

# PAC-Learning

## PAC-learning

- A concept class  $\mathcal{C}$  is said to be PAC-learnable if there exists an algorithm  $\mathcal{A}$  and a polynomial function  $poly(., ., ., .)$  such that for any  $\epsilon > 0$  and  $\delta > 0$ , for all distributions  $\mathcal{D}$  on  $\mathcal{X}$  and for any target concept  $c \in \mathcal{C}$ , the following holds for any sample size  $m \geq poly(1/\epsilon, 1/\delta, n, size(c))$

$$\mathbb{P}_{S \sim \mathcal{D}^m} [R(h_S) \leq \epsilon] \geq 1 - \delta$$

- if  $\mathcal{A}$  further runs in  $poly(1/\epsilon, 1/\delta, n, size(c))$ , the  $\mathcal{C}$  is considered to be efficiently PAC-learnable.
- When such  $\mathcal{A}$  exists, it is called a *PAC-learning algorithm* for  $\mathcal{C}$
- The parameter  $\delta > 0$  defines the confidence interval  $1 - \delta$  and  $\epsilon > 0$  the accuracy  $1 - \epsilon$ .

- Machine learning is fundamentally about **generalization**
- The problem comprises in selecting a function out of a *hypothesis set*, that is a subset of the family of all functions.
- The selected function is subsequently used to label all instances, including unseen examples
- How should a hypothesis set be chosen?
  - With a rich or complex hypothesis set, the learner may choose a predictor that is consistent with the training set
  - With a less complex one, it may have unavoidable errors on the training set
- Which one will lead to a better **generalization**?
- How should we define the complexity of a hypothesis set?




## What is generalization?

- *It is the ability of a model to adapt properly to **unseen data** drawn from the same distribution as the one used to create the model*
- Data are noisy, for different reasons
  - 1 errors during the acquisition phase
  - 2 errors in labeling the data points (i.e., teaching error)
  - 3 hidden or latent features
- We learn  $f$  by minimizing some variant of empirical risk, what can you say about the true risk?
- Two factors determine generalization ability:
  - 1 Model complexity
  - 2 Training set size


# Model Complexity & Generalization

# Understanding overfitting & underfitting




$$\beta_0 + \beta_1 x$$

Underfitting  $\mapsto$  “high bias”



$$\beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$$

“Just right”




$$\beta_0 + \beta_1 x + \dots + \beta_{15} x^{15}$$



“overfitting”  $\mapsto$  “high variance”

- In the **Overfitting** scenario, the learned hypothesis may fit the training set very well, but fail to **generalize** to new examples.
  - It is usually caused by a complicated function that creates various unnecessary curves and angles unrelated to the data
  - It has a large estimation error
- **Underfitting or high bias** occurs when the hypothesis function maps poorly to the trend of the data.
  - It is usually caused by a function that is very simple or that uses only few features
  - It has a large approximate error

## Generalization error vs. model complexity trade-off



# Fixed model complexity vs. dataset size



## Bias vs. variance trade-off

- **Bias** is the difference between the expected value of the estimator and the real value predicted by the estimator


$$\text{Bias}(f(x)) = \mathbb{E}[f(x) - y]$$

- A simple model has a high bias
  - High bias can lead to **underfitting**
- **Variance** is the deviation from the expected value of the estimates

$$\text{Var}(f(x)) = \mathbb{E}[(f(x) - \mathbb{E}(f(x)))^2]$$

- A complex model has a high variance
  - High variance usually leads to **overfitting**

# Bias vs. variance trade-off



## Addressing overfitting

- ① Reduce the number of features
  - Manually select which feature to keep
  - Model selection algorithm
- ② Regularization
  - Keeps all the features, but reduce the magnitude of the parameters
  - It works when there are many features contributing to predict  $y$



## Supervised learning assumption

- **Training set**  $\mathcal{D} = \{x^i, y^i\}_{i=1..n}$
- **Regression**  $y^i \in \mathbb{R}$
- **Classification**  $y^i \in \{0, 1\}$
- **Goal:** find a function  $f$  on the training set such that  $f(x^i) \approx y^i$
- **Empirical error** of  $f$  on the training set, given a loss function  $\mathcal{L}$

$$\mathbb{E}(f|\mathcal{D}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^i, f(x^i))$$

- **Regression**

$$\mathcal{L}(y^i, f(x^i)) = (y^i - f(x^i))^2$$

- **Classification**

$$\mathcal{L}(y^i, f(x^i)) = 1_{y^i \neq f(x^i)}$$

## Empirical error

- On the training set, it is a poor estimate of the ***generalization error***
- If the model is overfitting, the generalization error can be arbitrarily large
- Our goal is to estimate the generalization error on unseen data, which we might not have

# Validation & Cross-Validation

## Empirical vs. true risk

- In general, it is defined by

$$R(f) = R^{emp} + \text{overfit penalty}$$

- Overfit penalty depends on the complexity of the model
- **Regularization** approximates the overfit penalty. When the complexity of the model increases, we set up a larger overfit penalty
- **Cross-validation** tries to estimate  $R(f)$  directly

## Holdout method

- **Holdout method** is a popular approach for estimating the generalization performance of machine learning models
- Using **holdout method**, we split the initial dataset into training and test sets



- We want to choose a model that performs best on a **validation set** independent of the **training set**
- Since we have not used the validation data during the training phase, the validation set can be considered **unseen data**
- In this case, the error on the validation set is an estimation of the generalization error
- What is another issue in this approach?

## Model selection is a classification problem

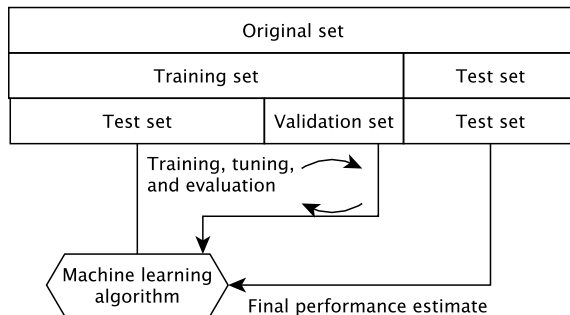
- We are interested in tuning and comparing different parameter settings to further improve the performance, for making prediction on unseen data
- This process is called ***model selection***
- Model selection refers to a given classification problem for which we want to select the optimal values of tuning parameters
- Therefore, if we reuse the same dataset over and over again during model selection, it will become part of our training data and thus the model will be more likely to overfit

## Dealing with multiple models

- What should we do if we want to choose among  $k$  different models?
  - 1 We have to train each model on the training set
  - 2 Then, compute the prediction error of each model on the validation set
  - 3 Finally, select the model with the smallest prediction error on the validation set
- In that case, what will be the generalization error?
  - It is hard to say
  - Validation data was used to select the model
  - Actually, as we have looked at the validation data, it is not anymore a good proxy for unseen data

## Holdout cross-validation

- A better way of using the holdout method for method selection comprises in splitting the dataset into three parts: a training set, a validation set, and a test set



- Therefore, the estimation error is sensitive to how we partition the training and the validation sets



## Handling the problem of validation set

- We have to set aside a test set that remains untouched during the training and the validation phases
- With the test set, we can use it to estimate the generalization error



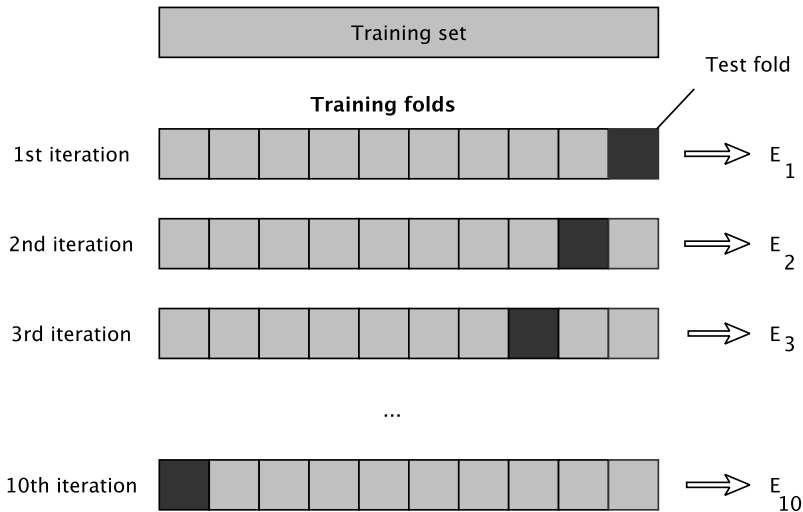
- How we decide the size of the training, validation, and test sets?
- How do we know that we have enough data to evaluate the prediction and the generalization errors?
- In **model selection**, we aim to pick the best model
- Whereas, in **model assessment**, we want to estimate the prediction errors on unseen data

We can use **cross-validation** and **bootstrap** techniques to empirically evaluate our model

## K-fold cross-validation

- ***k-fold cross-validation*** is a technique designed to give an accurate estimate of the true error without “wasting” too much data
- In the  $k$ -fold cross-validation, the original training set is partitioned into  $k$  folds without replacement
- $k - 1$  folds are used for the model training and one fold is used for testing
- For each fold, the model is estimated on the union of the other folds and then, the error of its output is estimated using the fold
- The average of all the errors is the estimate of the true error
- Once the best parameter is chosen, the model is retrained using the parameters of the entire training set

# K-fold cross-validation



## K-fold cross-validation algorithm

### Input:

training set  $S = (x^1, y^1), \dots, (x^p, y^p)$

set of parameter values  $\Theta$

learning algorithm  $\mathcal{A}$

$k$  (number of folds)

split  $S$  into  $S_1, S_2, \dots, S_k$

**foreach**  $\theta \in \Theta$  **do**

**for**  $i = 1..k$  **do**

$$h_{i,\theta} = \mathcal{A}(S \setminus S_i; \theta)$$

$$\text{error}(\theta) = \frac{1}{k} \sum_{i=1}^k \mathcal{L}_{S_i}(h_{i,\theta})$$

### Output:

$$\theta^* = \arg \min_{\theta} [\text{error}(\theta)]$$

$$h_{\theta^*} = \mathcal{A}(S; \theta^*)$$

## Cross-validation performance

- Estimating the prediction error

$$\begin{aligned} CV(f) &= \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^i, f_{k(i)}(x^i)) \\ &= \frac{1}{k} \sum_{l=1}^k \mathbb{E}(f|D_l) \end{aligned}$$

- where,  $f_{k(i)}$  is the  $k_i$ -th part of the data removed
  - $k_i$  is the fold in which  $i$  is
  - $D_l$  is the fold  $l$
- Estimating the expected prediction error

$$Error = \mathbb{E}[L(Y, f(X))]$$

## Cross-validation issues

- The training set becomes  $(k - 1) * n/k$ 
  - small training set may lead to biased estimator of the error
- A special case of the k-fold cross-validation is the **leave-one-out (LDO)**; i.e.,  $k = n$ 
  - approximately unbiased of the expected prediction error
  - potential high variance, since the training sets are similar to each other
  - computation can be very difficult
- In practice,  $k$  is set up to 5 or 10.

# Bootstrap

- Randomly draws datasets with replacement from the training set
- Repeats  **$B$**  times (often,  $B = 100$ ), which leads to  **$B$**  models
- Leave-one-out bootstrap error
  - for each training point  $i$ , predicts with the  $b_i < B$  models that did not have  $i$  in their training set
  - computes the average prediction errors
- This leads for training set that has  $0.632 * n$  distinct examples. Why?

$$\begin{aligned}\mathbb{P}(i \in x_k) &= 1 - \left(1 - \frac{1}{n}\right)^n \\ &\approx 1 - e^{-1} \\ &= 0.632\end{aligned}$$

- It has a high computational cost

# Assessing Model Performance



## Model evaluation metrics

- **Precision, recall, and F1-score** are performance metrics that can be used to measure a model's relevance
- The performance of a model can be summarized by means of a **confusion matrix**

		Predicted class	
		+	-
Actual class	+	True Positives (TP)	False Negatives (FN)
	-	False Positives (FP)	True Negatives (TN)

- Each row refers to actual classes recorded in the test set, and each column to classes as predicted by the predictor
- **False positives** represent **false alarms**, which are also known as **type I errors**
- **False negatives** represent **misses classifications**, which are called **type II errors**

## Computing precision, recall, and F1-score

- Prediction **error** (**ERR**) and **accuracy** (**ACC**) provide general information about how many samples are misclassified
- **Error (ERR)**

$$ERR = \frac{FP + FN}{FP + FN + TP + TN}$$

- **Accuracy (ACC)**

$$ACC = \frac{TP + TN}{FP + FN + TP + TN} = 1 - ERR$$

- **Sensitivity = Recall = True Positive Rate (TPR)**

$$TPR = \frac{TP}{FN + TP}$$

## Computing precision, recall, and F1-score

- **False Positive Rate (FPR)**

$$FPR = \frac{FP}{FP + TN}$$

- **Specificity = True Negative Rate (TNR)**

$$TNR = \frac{TN}{FP + TN}$$

- **Precision = Positive Predictive Value (PPV)**

$$Precision = \frac{TP}{TP + FP}$$


- **F1-score** represents the harmonic mean of precision and sensitivity

$$F1 = \frac{2TP}{2TP + FP + FN}$$

## ROC curves

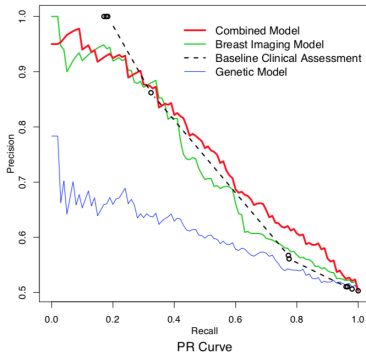

- **Receiver operator characteristic (ROC)** is a tool for selecting models for classification based on their performance with respect to the **false positive** and **true positive** rates
- The diagonal of an ROC plot can be interpreted as a random guessing
- It is summarized by the **area under the curve (AUC)**, which characterize the performance of a classification model

# ROC curves



- **Perfect classifier:**  
 $AUC = 1.0$
- **Random classifier:**  
 $AUC = 0.5$
- **Our classifier:**  
 $0.5 < AUC < 1.0$

## Example: Breast Cancer Risk Prediction on Mammograms



Predicting breast cancer risk based on mammography images. **Source:** Liu et al. (2013)<sup>1</sup>

- **High recall** means less chances to miss a case
- **High precision** means substantially more true diagnoses than false alarms

---

<sup>1</sup> Jie Liu et al. "Genetic variants improve breast cancer risk prediction on mammograms". In: *Annual Symposium Proceedings*. Vol. 2013. 2013, p. 876.

## Assessing regression model performance

- ***Residual sum of squares (RSS)***

$$RSS = \sum_{i=1}^n (y_i - f(x^i))^2$$

- ***Root-mean squared error (RMSE)***

$$RMSE = \sqrt{\frac{\sum_{i=1}^n (y_i - f(x^i))^2}{n}}$$

## Assessing regression model performance

- ***Relative squared error (RSE)***

$$RSE = \frac{\sum_{i=1}^n (y_i - f(x^i))^2}{\sum_{i=1}^n (y^i - \bar{y})^2}$$

- ***Coefficient of determination***

$$R^2 = 1 - RSE$$



## References

- Marianthi Markatou et al. “Analysis of Variance of Cross-Validation Estimators of the Generalization Error”. In: *Journal of Machine Learning Research* 6 (2005), pp. 1127–1168
- Bradley Efron and Robert Tibshirani. “Improvements on cross-validation: the 632+ bootstrap method”. In: *Journal of the American Statistical Association* 92.438 (1997), pp. 548–560
- L. G. Valiant. “A Theory of the Learnable”. In: *Communication of the ACM* 27.11 (1984), pp. 1134–1142
- Hal Daume III. *A Course in Machine Learning*. 2nd. Self-published, 2017. URL:  
[http://ciml.info/dl/v0\\_99/ciml-v0\\_99-all.pdf](http://ciml.info/dl/v0_99/ciml-v0_99-all.pdf)
  - 1 **Noise**: session 2.3
  - 2 **Overfitting**: session 2.4
  - 3 **Bias-variance trade-off**: session 5.9
  - 4 **Holdout method**: session 2.5
  - 5 **Cross-validation**: session 5.6

## References

- ⑥ **Assessing model performance:** session 5.5
- Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. 2nd. Springer, 2016. URL:  
<https://web.stanford.edu/~hastie/Papers/ESLII.pdf>
- ① **Overfitting:** session 7.1
- ② **Bias-variance trade-off:** sessions 2.9, 7.2, and 7.3
- ③ **Cross-validation:** session 7.10
- ④ **Bootstrap:** session 7.11